

Thread 3 — Chromatogram Integration: allocation analysis

Why it is not a quick allocation win, unlike Quant & Output

2026-06-24 | branch `perf/runtime-profiling`

Bottom line. The Chromatogram Integration phase (~3.4 GB cumulative allocation, warm 2-file MTAC) is dominated by the *actual chromatogram signal* being materialized, plus one structure-of-arrays conversion copy. Unlike Quant & Output — whose 3.9 GB win came from skipping *redundant* string churn (99.5% duplicate accessions) — Integration’s allocation is mostly *necessary numeric data* that cannot be skipped. The single reducible item (the conversion copy) needs a moderate refactor, so it does *not* clear the “no excessive complexity” bar as a quick win. Recommendation: measure the copy precisely before deciding.

1 What the phase does

Three stages, in order:

1. `build_chromatograms` (`utils.jl:707`) — walks the scan range, runs the fused deconvolution kernel per scan, and **emits one** `MS2ChromObject` **per (precursor, scan) chromatogram point** into a `Vector{MS2ChromObject}` (pre-sized to `n_scans`×100). It returns `DataFrame{@view(chromatograms[utils.jl:878])}`.
2. `integrate_precursors` (`utils.jl:399`) — groups the chromatogram table by precursor and integrates each peak.
3. `integrate_chrom` (`integrate_chrom.jl`) — per precursor: Whittaker–Henderson smoothing → second-derivative bounds → baseline-subtracted trapezoidal area.

2 Where the 3.4 GB is — and is not

The integration loop is already allocation-lean. `integrate_precursors` (stages 2–3) was written to reuse memory:

- per-thread `WHWorkspace` and `Chromatogram` state are pre-allocated once (`utils.jl:431--432`) and reused across all precursors;
- every chromatogram slice is passed as an `@view` (`utils.jl:464--467`) — no per-precursor copies;
- results are written into caller-provided output vectors (`peak_area`, `new_best_scan`, `points_integrated`).

So the per-precursor smoothing/integration is *not* the allocator. (The line profiler agreed: `WHSmooth!` ≈0.11 GB, `build_chromatograms` leaf ≈0.15 GB — small.)

The allocation lives in stage 1, and it is mostly real data:

- **the `Vector{MS2ChromObject}` of chromatogram points** — this *is* the signal being integrated (every (precursor, scan) intensity in each precursor’s RT window). It scales with the data, not with waste.
- **the `AoS→DataFrame` conversion copy** (`DataFrame(@view(...))`, `utils.jl:878`) — the array-of-structs `MS2ChromObject` vector is copied field-by-field into the `DataFrame`’s columns; the original vector is then discarded. This is a transient $\sim 1\times$ duplicate of the chromatogram data, analogous to the Main Search `df_materialize` (0.78 GB there).

3 Why this is unlike the Quant & Output wins

	Quant & Output (Thread 2)	Chrom Integration (Thread 3)
What allocates	per-row <code>split</code> + <code>String</code> interning	chromatogram signal + <code>AoS→DF</code> copy
Nature	redundant (99.5% duplicate rows)	necessary (numeric data)
Fix	memoize & skip repeats	emit columns (SoA) instead of copy
Complexity	low (a <code>Set</code> /cache)	moderate (restructure emit path)
Result	−92% / −55%, bit-identical	only the copy is reducible

Thread 2 worked because the allocation was *repeated work over duplicate strings* — skippable with a tiny memo. Thread 3 has no such redundancy: the chromatogram points are distinct measured intensities. The only thing we are paying *twice* for is the `AoS` vector plus its `DataFrame` copy.

4 The one available lever, and the recommendation

Avoiding the conversion copy means having `build_chromatograms` write directly into column vectors (`rt`, `scan_idx`, `intensity`, `precursor_fraction_transmitted`, ...) — a structure-of-arrays emit path — and constructing the `DataFrame` from those columns with no per-element copy. That removes the duplicate but is a **moderate refactor** of the emit hot loop (and must stay bit-identical and not regress speed).

Recommendation. Before refactoring, *measure*: add `gc_bytes` sub-attribution (same method as the Main Search `MSLOOP` pass) to split the 3.4 GB into (a) the chromatogram `Vector`, (b) the `DataFrame` conversion copy, and (c) everything else. Refactor *only* if the copy is a clean $\gtrsim 1$ GB and the SoA emit stays simple. Otherwise, conclude Thread 3: the phase is near its necessary-data floor and further cuts would trade real complexity for little memory.